

Concurrency Control

7

The system must control the interaction among the concurrent executing transactions. This control is achieved through one of Concurrency-control schemes.

Lock-Based protocols:

- one way to ensure serializability is to require that data items be accessed in a mutually exclusive manner.
↓
if T is accessing a data item, no other T can modify that data item.
- allow a T to access a data item only if it is currently holding a lock on that item.

Locks:

VARIOUS

Data items can be locked in two modes:

- ① Shared (S) mode: Data item can only be read.
If a Transaction T_i has obtained a shared-mode lock on item a_1 , then T_i can read, but cannot write, a_1 .
- ② Exclusive (X) mode: Data item can be both read as well as written. If a ~~T_i~~ has obtained a exclusive-mode lock on item a_1 , then T_i can both read and write a_1 .

- we require that every transaction request a lock in an appropriate mode on data item a_1 , depending on the types of operations that it will perform on a_1 .
- the transaction makes the request to the concurrency-control committee resolution manager.
- the T can proceed with the transaction resolution after getting.

→ compatibility function:

A $\xrightarrow{\text{B}}$ makes at lock.

	S	X
S	True	F
X	*	F

Fig: Lock-compatibility matrix group

→ If T_i request a lock of mode A on item on, on which T_j currently holds mode B.

→ If T_i can be granted a lock on immediately, even the presence of the mode B lock, then mode A is compatible with mode B.

In the above matrix, shared mode is compatible with shared mode.

→ Several shared-mode locks can be held simultaneously on a same data item.

→ A subsequent exclusive-mode lock request has to wait until the currently held shared-mode locks are released.

T_1 : lock-X(B);
read(B);
 $B := B - \$0$;
write(B);
unlock(B);
lock-X(A);
read(A);
 $A := A + \$0$;
write(A);
unlock(A)

T_2 : lock-S(A);
read(A);
unlock(A);
lock-S(B);
read(B);
unlock-B(B);
display(A+B)

T_1	T_2	concurrently-control mgt
lock-X(B)		grant-X(B, T_1)
read(B)		
$B := B - \$0$		
write(B)		
unlock(B)		
	lock-S(A)	grant-S(A, T_2)
	read(A)	
	unlock(A)	
	lock-S(B)	grant-S(B, T_2)
	read(B)	
	unlock(B)	
	display(A+B)	grant-X(A, T_1)

access a data item, T_i must first lock that item.
 If the data item is already locked by T_j in incompatible mode, the Concurrency-control mgr will not grant the lock until all incompatible locks held by other transactions have been released.

~~So T_i~~

- a 'T' must hold a lock on a data item as long as it accesses that item.
- for a 'T', to unlock a data item immediately after its final access of that data item is not always desirable, since serializability may not be ensured

$A = 100 \$ \quad B = 200 \$$

$\text{A} / \text{B} - T_2 = 300 \$ \quad \text{If } \cancel{\text{A}} =$

In the above case

T_2 displays $200 \$$ which is incorrect.

$T_3:$
 $\text{lock-X}(B);$
 $\text{dead}(B);$
 $B := B - 50;$
 $\text{write}(B);$
 $\text{lock-X}(A);$
 $\text{read}(A);$
 $A := A + 50;$
 $\text{write}(A);$
 $\text{unlock}(B);$
 $\text{unlock}(A)$

$T_4:$
 $\text{lock-S}(A).$ db is correct with
 $\text{dead}(A);$
 $\text{lock-X}(B);$
 $\text{dead}(B).$
 $\text{display}(A+B);$
 $\text{unlock}(A);$
 $\text{unlock}(B);$

$T_3 \leftarrow T_4.$

~~Consi~~
 → executing $\text{lock-S}(B)$ causes T_4 to wait for T_3 to release its lock on B, while executing $\text{lock-X}(A)$ causes T_3 to wait for T_4 to release its lock on A.

→ neither of these transactions can ever proceed with its normal execution.

This situation is called a deadlock.

→ In handle deadlock ...

T_3	T_4
$\text{lock-X}(B)$	
$\text{dead}(B);$	
$B := B - 50$	
$\text{write}(B)$	$\text{lock-S}(A)$
	$\text{dead}(A)$
	$\text{lock-S}(B)$
	$(\text{lock-X}(A))$

fig: which leads to deadlock.

A locking protocol is a set of rules followed by all transactions while requesting and releasing locks.

→ Locking protocols restrict the set of possible schedules.

Starvation:

$T_2 \rightarrow$ has a shared mode lock on data item

T_1 requests an X-mode lock on that data item

T_1 has to wait until T_2 releases the lock.

In the meanwhile, T_3 may request a S-mode lock on same data

item. T_2 grants the S-mode lock.

T_2 may release the lock, but still T_1 has to wait for T_3 to finish

so T_1 may never make progress, and is said to be starved.

The Two-phase locking protocol:

→ This is a protocol which ensures Conflict-Serializable Schedules.

Phase 1: Growing phase:

- transaction may obtain locks.

- " may not release any lock.

Phase 2: Shrinking phase:

- transaction may release locks.

- " may not obtain any new locks.

T_3 & T_4 are two phase, T_1 & T_2 are not two phase

→ ~~if~~ this protocol ensures Conflict Serializability.

It can be proved that the transactions can be serialized in the order of their lock points (i.e. the point where a transaction acquired its final lock).

~~Two phase locking does not ensure freedom from deadlocks.~~

Cascading roll backs is possible under two-phase locking.

To avoid this, follow a modified protocol called

Strict two-phase locking protocol: a transaction locking be two phase, & must hold all its exclusive locks till it commits / aborts.

Rigorous two-phase locking protocol: here all locks are held till commit / abort. In this protocol transactions can be serialized in the order in which they commit.

T_g: read(a₁);
read(a₂);
= = =
read(a_n);
write(a₁).

T_g: read(a₁)
read(a₂)
display(a₁, a₂).

→ Scheduling 2 to 2 ~~two-phase~~

Two-phase locking with lock conversions.

→ Convert a lock-S to a lock-X
is upgrading

→ Convert a lock-X to a lock-S is
downgrading.

→ Upgrading takes place in only the
growing phase

→ Downgrading takes place in only the
shrinkage phase.

→ This protocol ensures serializability.
& the schedules are cascadeless.

T _g	T _g	T _g	T _g
lock-S(a ₁)			lock-S(a ₁)
lock-S(a ₂)			lock-S(a ₂)
lock-S(a ₃)			
lock-S(a ₄)			
			unlock(a ₁)
			unlock(a ₂)
lock-S(a _n)			
upgrade(a)			
			write(a)
			unlock -

Automatic acquisition of locks.

- when a transaction T_i issues a $\text{read}(o_i)$ operation, the system issues a lock- $S(o_i)$.
- when T_i issues a $\text{write}(o_i)$ operation, the system checks to see whether T_i already holds a shared lock on o_i , if it does, then the system issues an upgrade (o_i) instl., followed by the write(o_i) instl., otherwise, the system issues a lock- $X(o_i)$ instl/
- All locks obtained by a transaction are unlocked after Commit or aborts.

Graph-Based protocols:

- It is an alternative to two-phase locking ..
- Impose a partial ordering \rightarrow on the set $D = \{d_1, d_2, \dots, d_n\}$ of all data items.
- If $d_i \rightarrow d_j$ then any transaction accessing both d_i & d_j , must access d_i before accessing d_j -
- Partial ordering implies that the set D may now be viewed as a directed acyclic graph, called a database graph.
- Tree-protocol: is a simple kind of graph protocol, which employs only exclusive locks
- only exclusive locks are allowed.
- each T_i must lock a data item at most once.
→ rules are:
 - ① The first lock by T_i may be on any data item
 - ② subsequently, a data item o_i can be locked by T_i only if the parent of o_i is currently locked by T_i .
 - ③ Data items may be unlocked at any time.
 - ④ A data item that has been locked . . .

Time-Stamp based protocols:

Time stamps: Time stamp is a unique identifier created by DBMS to identify a transaction T_i in the system, a unique fixed timestamp is assigned as $TS(T_i)$ → before the transaction T_i starts execution.

If T_i has been assigned timestamp $TS(T_i)$ and a new transaction T_j enters the system, then $TS(T_i) < TS(T_j)$.
→ use the value of system clock as the timestamp.

Timestamp of T_i is equal to the value of clock when the transaction enters the system.

If $TS(T_i) < TS(T_j)$, then the produced schedule is equivalent to a serial schedule in which T_i appears before T_j for each data item a , two timestamp values are assigned.

→ w-timeStamp(a): denotes the largest (latest) timestamp of any transaction that executed $\text{write}(a)$ successfully.

→ R-timeStamp(a): denotes the largest (latest) timestamp of any transaction that executed $\text{read}(a)$ successfully.

→ these timestamps are updated whenever a new $\text{read}(a)$ or $\text{write}(a)$ instruction is executed.

The time-stamp-ordering protocol:

→ It ensures that any conflicting read and write operations are executed in timestamp order.

The protocol operates as follows

① Suppose T_i issues $\text{read}(a)$.

② If $TS(T_i) < w\text{-timestamp}(a)$, then T_i needs to read

T_{10}	T_{11}	T_{12}	T_{13}
$\text{lock}_X(B)$	$\text{lock}_X(D)$ $\text{lock}_X(H)$ $\text{unlock}(D)$		
$\text{lock}_X(E)$			
$\text{lock}_X(D)$			
$\text{unlock}(B)$			
$\text{unlock}(E)$			
	$\text{unlock}(H)$		
		$\text{lock}_X(B)$ $\text{lock}_X(F)$	
unlock			
$\text{lock}_X(G)$ $\text{unlock}(D)$			
		$\text{lock}_X(D)$ $\text{lock}_X(H)$ $\text{unlock}(D)$ $\text{unlock}(F)$ $\text{unlock}(H)$	
$\text{unlock}(G)$		$\text{unlock}(B)$	

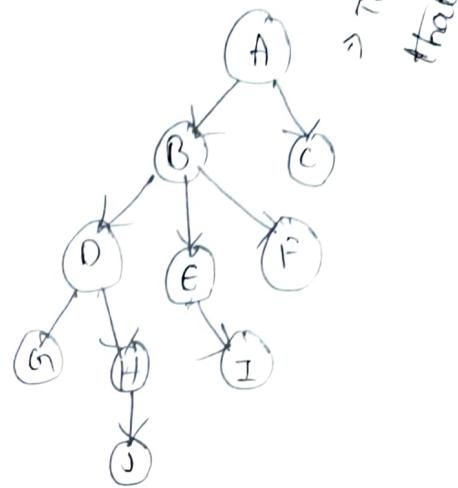


fig. Tree-structured Database graph.

fig. Serializable schedule under the tree protocol.
 → all schedules are conflict serializable schedules according to graph boundaries.
 → Tree protocol does not ensure recoverability and cascadelness
 → To ensure that one, the protocol can be modified to not permit releases of locks until the end of the transaction
 → holding locks improves concurrency & ensures recoverability.

adv/ → It is a dead-lock free, so no rollbacks are required.
 → unlocking may occur earlier. So it leads to shorter waiting time, & to an increase in concurrency.

drawback → A transaction may have to lock data items that it doesn't access. It increases the waiting time, & decrease in concurrency.

- The protocol generates schedules that are not recoverable.
- To ensure recoverability
- Recoverability & cascadelessness
Can be ensured by performing all writes together at the end of the transaction.
- It can be guaranteed by using a limited form of locking.

T ₁₄	T ₁₅
read(B)	
	B = B - SO
	write(B)
read(A)	
	A = A + SO
	write(A)
	display(A+B)

Thomas' write Rule:

- modification to the TS-81 protocol, that allows greater concurrency if $TS(T_6) < TS(T_7)$
So above schedule is not tsb in TS-OP

T ₁₆	T ₁₇
read(B)	
write(B)	write(A)

The modification to the TS-ordering protocol, called Thomas' write rule rule. If T_i issues a $write(g)$.

- ① if $TS(T_i) < R\text{-timestamp}(g)$, the system rejects write operation & T_i rolls back.
- ② if $TS(T_i) \leq w\text{-timestamp}(g)$, hence this write operation can be ignored.
- ③ otherwise, the system executes the write operation & $w \rightarrow TS(g) \neq TS$

(b) If $TS(T_i) \geq w\text{-timestamp}(a)$, then read operation is exec.
& $R\text{-}TS(a)$ is set to $TS(T_i)$.

(2) Suppose T_i issues $w\text{-write}(a)$

(a) If $TS(T_i) < R\text{-Timestamp}(a)$, the value of a in that T_i .
System rejects write operation and T_i rolls back.

(b) If $TS(T_i) \geq w\text{-timestamp}(a)$.
System rejects this write operation & T_i rolls back.

(c) otherwise, the system executes the write operation and
set $w\text{-TS}(a)$ to $TS(T_i)$.

If a Transaction T_i is rolled back, the system assigns
a new timestamp and restarts it's.

T_{14} : $\text{read}(B);$
 $\text{read}(A);$
 $\text{display}(A+B);$

$$TS(T_{14}) < TS(T_{15})$$

& the schedule is possible under the
time-stamp protocol.

T_{15} : $\text{read}(B)$
 $B = B - 50;$
 $\text{write}(B);$
 $\text{read}(A);$
 $A = A + 50;$
 $\text{write}(A);$
 $\text{display}(A+B).$

Some schedules are
not possible under the
two-phase locking protocol
but timestamp protocol.

- ~~It~~ this TS-ordering protocol ensures conflict serializability.
because conflicting operations are processed in timestamp order.
- It ensures freedom from deadlock, so no transaction ever waits.
- But ability of starvation.

Validation-Based protocols:

If many of the T's are read-only transactions, the rate of conflicts among transactions may be low, without supervision of concurrency-control. In fact, it leads db to inconsistent state. But under the supervision, it may impose overhead. To impose less overhead, you can use these schemes.

T_i executes in two or three different phases.

① Read phase: During this phase, the system executes T_i . It reads data items & stores in local variable.

It performs write operation on temporary local variables, ~~but~~ without update of the actual db.

② Validation phase:

T_i performs a validation test to determine whether it can copy to the db the temporary local variables that had the results of write operations without causing a violation of serializability.

③ Write phase: If T_i succeeds in validation, then the system applies the actual updates to the db. Otherwise the system rolls back.

Three different time stamps with T_i

④ Start(T_i), the time when T_i started its execution.

⑤ Validation(T_i), the time when T_i finished its read phase and started its validation phase.

⑥ Finish(T_i), the time when T_i finished its write phase.

The value

If $TS(T_j) < TS(T_k)$, then any produced schedule must be equivalent to a serial schedule.

But we have choose validation(T_i) rather than start(T_i).

The validation test for transaction T_i requires that, for all transaction T_j with $TS(T_i) < TS(T_j)$, one of the following two conditions must hold.

- ① $Finish(T_i) < Start(T_j)$, T_i completes its execution before T_j starts,
- ② The set of data items written by T_i does not interact with the set of data items read by T_j , and T_i completes its write phase before T_j starts its validation phase.
i.e. $Start(T_j) < Finish(T_i) < validation(T_j)$.

writes of T_i & T_j does not overlap.

Validation phase succeeds.

This schedule is serializable.

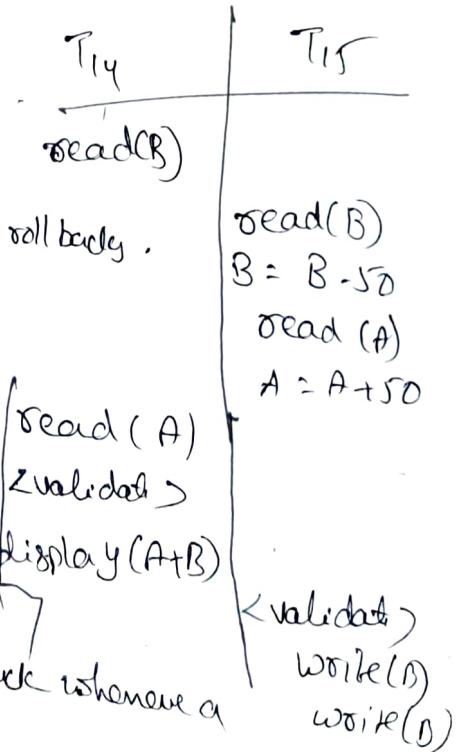
→ Validation automatically avoids cascading roll back.

→ Possibility of starvation.

→ Validation Scheme is called the

optimistic concurrency - control scheme
they will be able to finish execution & validate display ($A+B$)
& locking & timestamp are pessimistic

In that they use force writing a roll back whenever a conflict is detected, ...



deadlock handling:

- A system is in deadlock state if there exists a set of transactions such that for another transaction in the set is waiting for another transaction in the set.
- There exists a set of waiting transaction $\{T_0, T_1, \dots, T_n\}$ such that T_0 is waiting for data item that T_1 holds, & T_1 is waiting for data item that T_2 holds, and ..., and T_{n-1} is waiting for " " " " T_n holds, and T_n is waiting for " " " " T_0 holds, none of the transactions can make progress in such a situation.
- Rolling back some or all is the remedy.
- Deadlock prevention protocol to ensure that the system will never enter a deadlock state.
- deadlock detection and deadlock recovery scheme.

We can allow the system to enter a deadlock state, and then try to recover by ~~breaking~~.

Deadlock prevention:

- 1st approach → ensures that no cyclic waits can occur by ordering the requests for locks, & requiring all locks to be acquired together.
- Perform transaction rollback instead of waiting for a lock.
- 1st approach → each transaction locks all its data items before it begins execution.
 - ~~disadv~~ → 1) it is often hard to predict, before the transaction begins, what data items need to be locked
 - 2) data-item utilization

Preemption and Transaction rollbacks.

In preemption, when a transaction T_2 requests a lock on a transaction T_1 , which holds, the lock granted to T_1 may be preempted by rolling back of T_1 , and granting of the lock to T_2 .

→ timestamp is assigned to each transaction.

two different deadlock prevention schemes using TS's.

1) wait-die:

is a non-preemptive technique.

when T_j , T_i is allowed to wait only if it has a timestamp smaller than that of T_j . otherwise T_i is rolled back (die).

T_{22} requests a data item held by T_{23} .

T_{22} has to wait.

$T_{22}:5 \quad T_{23}:10 \quad T_{24}:15$

~~T_{22} requests a data item held by T_{23} .~~

T_{24} requests a data item held by T_{23} .

~~T_{24} will be rolled back.~~

2) wound-wait: is a preemptive technique.

when T_i requests a data item currently held by T_j , and T_i is allowed to wait only if it has a timestamp larger than that of T_j . otherwise, T_i is rolled back & preempted.

→ whenever the system rolls back T_i 's, it is impl. to ensure that there is no starvation.. i.e. no transaction gets rolled back repeatedly. Both schemes avoid starvation.

- each node in the tree can be locked individually.
- when a 'T' locks a node in 'S' or 'X' mode, all the descendants are ^{implicitly} locked in the same mode.
- Ti gets an explicit lock on file F_c in exclusive mode, then it has an implicit lock in exclusive mode on all the records belonging to that file. no need to lock the individual records of F_c explicitly.
- intention lock modes
 - If a node is locked in an intention mode, explicit locking is being done ~~at~~ at lower levels of the tree.
- Intention locks are put on all the ancestors of a node before that node is locked explicitly.
-

Multiple Granularity

- grouping of several data items, & treat them as one \rightarrow to individual unit. If T_i needs to access the entire db, & a locking protocol is used then T_i must lock each item in the db.
- It is better to make a single lock request to lock the entire db.
- allows the system to define multiple levels of granularity, allowing data items to be of various sizes and defining a hierarchy of data. Data are nested within one another. Such a hierarchy can be represented graphically as a tree.
- nonleaf-node represents the data associated with its descendants.
- Each node is an independent data item.

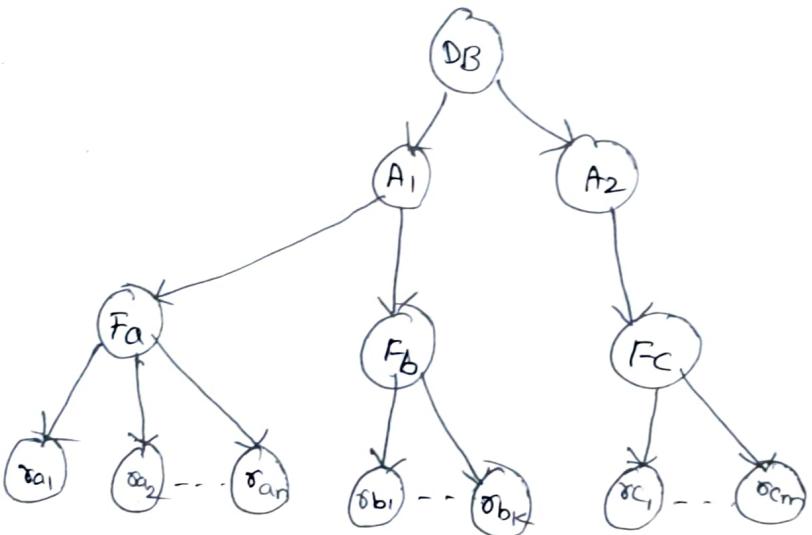


Fig: Granularity hierarchy.

- It consists of four levels of nodes.
- highest level represents the entire db.
- below it are nodes of type area (db consists of exactly these areas).
- file - each area contains exactly those files that are its child nodes.
- Each file has nodes of type record.

Deadlock Detection and Recovery:

- The system must
 - Maintain info. about the current allocation of data items to transactions, as well as any outstanding data item requests.
 - provide an algol. that uses this info. to determine whether the system has entered a deadlock state.
 - Recover from the deadlock when the detection algol. determines that a deadlock exists

Deadlock Detection:

Deadlocks can be described in terms of directed graph called a wait-for-graph.

When T_i requests a data item currently being held by T_j , then the edge $T_i \rightarrow T_j$ is inserted in the wait-for-graph.

→ A deadlock exists in the system if and only if the wait-for-graph contains a cycle. Each transaction involved in the cycle is said to be deadlocked. To detect deadlocks, the system needs to maintain the wait-for-graph and periodically to invoke an algol. that searches for a cycle in the graph.

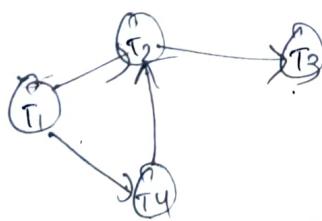


fig: wait-for-graph with no cycle

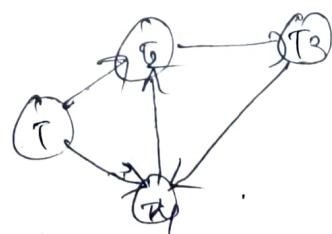


fig: wait-for-graph with a cycle

Recd

Recovery from Deadlock:

- when a detection algo. determines that a deadlock has occurred, the system must recover from the deadlock.
- i.e. roll back one or more transaction to break the deadlock.
- 1) Selection of a victim: we must determine which transaction to roll back to break the deadlock.
 - 2) Roll back: once we have decided that a particular transaction must be rolled back, we must determine how far this transaction should be rolled back.
 - 1) total rollback - i.e. Abort the transaction & restart it.
 - 2) partial rollback - which looks the selected transaction ready to release in order to break the deadlock.
 - 3) Starvation: it may happen that the same transaction is always picked as a victim. Thus this transaction never completes its task, so there is a starvation. We must ensure that a transaction can be picked as a victim only a finite no. of times.

Deferred Database modification:

- It defers all modifications to the log, but deferring the execution of all work operations of a transaction after partially commits.
- Transactions execute serially.
- When a Transaction partially commits, the info. entire log associated with the transaction is used in executing the deferred writes.
- Transaction starts by writing $\langle T_i \text{ start} \rangle$ record to log
- A write(x) operation results in a log record $\langle T_i, x, v \rangle$ being written, V is new value
the work is not performed on x at this time, it is delayed
- When T_i partially commits, $\langle T_i \text{ commit} \rangle$ is written to the log.
- Finally, the log records are read and used to actually execute the previously deferred writes.

To : read(A)
 $A := A - 50$
write(A)
read(B)
 $B := B + 50$
write(B)

T_1 : read(c)
 $C := C - 100$
write(c)

To follow by T_j ,

- During recovery after a crash, a transaction needs to be redone if and only if both $\langle T_i \text{ start} \rangle$ and $\langle T_i \text{ commit} \rangle$ are there in the log.
- Redoing a transaction T_i (redo T_i) - sets the value of all data items updated by the transaction to the new values

Log-Based Recovery:

(15)

- log is used for recording database modifications.
- the log is a sequence of log records, recording all the update activity in the db.
- An update log record describes a single database work.
It has 4 fields:
 - ① Transaction identifier: is the unique identifier of the transaction that performed the write operation.
 - ② Data-item identifier: " " " " data item written
 - ③ Old value: is the value of the data item prior to the write
 - ④ New value: is the value that the data item will have after the write.Various types of log records are.

$\langle T_i \text{ start} \rangle$. T_i has started.

$\langle T_i, X_j, v_1, v_2 \rangle$. T_i performed a write on data item X_j ,
 X_j had value v_1 , before the write & will have
Value v_2 after the write.

$\langle T_i \text{ commit} \rangle$. T_i has committed.

$\langle T_i \text{ abort} \rangle$. T_i has aborted.

⇒ whenever a transaction performs a write, it is essential
that the log record for that write be created before the
db is modified.

- once log record exists, we apply the modification to the db.
- undo a modification that has already been applied to the db.
- log records are written directly to stable storage,
- log contains a complete record of all db activity.

Two approaches using logs

→ Deferred database modification